

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3
        4 df = pd.read_csv("exchange-rates-2000-2022.csv")
        5 df
```

```
Out[1]:
```

	DATE	USD	GBP	EUR	JPY100	CHF	AUD	CAD	SGD	HKD100
0	3-Jan-00	3.8	6.1598	3.8646	3.7259	2.4066	2.5002	2.6288	2.2885	48.8715
1	4-Jan-00	3.8	6.2113	3.8956	3.7079	2.4262	2.5021	2.6257	2.2923	48.863
2	5-Jan-00	3.8	6.2246	3.911	3.6936	2.4382	2.4915	2.615	2.2945	48.8579
3	6-Jan-00	3.8	6.2466	3.9239	3.6442	2.4453	2.5002	2.6206	2.2948	48.8586
4	7-Jan-00	3.8	6.2603	3.9186	3.6087	2.4417	2.4854	2.6043	2.2843	48.8535
...
17012	23-Dec-22	0.0187	0.1073	2.78	0.2115	5.3486	120.569	1.9574	3.343	0.1789
17013	27-Dec-22	0.0188	0.1072	2.7768	0.2112	5.3442	120.3763	1.953	3.3403	0.1786
17014	28-Dec-22	0.0188	0.1074	2.7789	0.2117	5.3502	120.6643	1.9557	3.3439	0.1794
17015	29-Dec-22	0.0187	0.107	2.799	0.2111	5.3399	120.3169	1.9488	3.3374	0.1784
17016	30-Dec-22	0.0187	0.1069	2.7899	0.2108	5.3342	120.1601	1.9466	3.3338	0.1783

```
In [2]: 1 #A. Explanatory Data Analysis (EDA)
```

```
In [3]: 1 df.shape
```

```
Out[3]: (17017, 10)
```

```
In [4]: 1 df.columns
```

```
Out[4]: Index(['DATE', 'USD', 'GBP', 'EUR', 'JPY100', 'CHF', 'AUD', 'CAD', 'SGD',
              'HKD100'],
              dtype='object')
```

```
In [5]: 1 df.columns = map(str.upper, df)
        2 df.rename(columns=lambda x:x+'_MYR', inplace=True)
        3 df.rename(columns={'DATE_MYR':'Date'}, inplace=True)
```

In [6]: 1 df.head(15)

Out[6]:

	Date	USD_MYR	GBP_MYR	EUR_MYR	JPY100_MYR	CHF_MYR	AUD_MYR	CAD_MYR	SGD_
0	3-Jan-00	3.8	6.1598	3.8646	3.7259	2.4066	2.5002	2.6288	2
1	4-Jan-00	3.8	6.2113	3.8956	3.7079	2.4262	2.5021	2.6257	2
2	5-Jan-00	3.8	6.2246	3.911	3.6936	2.4382	2.4915	2.615	2
3	6-Jan-00	3.8	6.2466	3.9239	3.6442	2.4453	2.5002	2.6206	2
4	7-Jan-00	3.8	6.2603	3.9186	3.6087	2.4417	2.4854	2.6043	2
5	11-Jan-00	3.8	6.2227	3.9007	3.6132	2.4225	2.4909	2.6107	2
6	12-Jan-00	3.8	6.2641	3.9265	3.588	2.4381	2.5017	2.6088	2
7	13-Jan-00	3.8	6.2569	3.9136	3.6002	2.4294	2.5175	2.6146	2
8	14-Jan-00	3.8	6.2554	3.8999	3.5798	2.4184	2.5314	2.6201	2
9	17-Jan-00	3.8	6.2183	3.8509	3.6087	2.3868	2.528	2.6185	2
10	18-Jan-00	3.8	6.2094	3.8384	3.6201	2.378	2.5211	2.623	2
11	19-Jan-00	3.8	6.2238	3.8559	3.5978	2.3904	2.5245	2.6195	
12	20-Jan-00	3.8	6.2468	3.8429	3.6081	2.3813	2.5207	2.6166	2
13	21-Jan-00	3.8	6.2805	3.8623	3.6208	2.3967	2.5291	2.6282	2
14	24-Jan-00	3.8	6.2654	3.8181	3.6213	2.3676	2.4972	2.6335	2

```
In [7]: 1 df.isnull().sum()
```

```
Out[7]: Date                2
        USD_MYR             2
        GBP_MYR             2
        EUR_MYR             2
        JPY100_MYR          2
        CHF_MYR             2
        AUD_MYR             2
        CAD_MYR             2
        SGD_MYR             2
        HKD100_MYR          2
        dtype: int64
```

```
In [8]: 1 df = df.dropna().reset_index(drop=True)
        2 df.isna().sum()
```

```
Out[8]: Date                0
        USD_MYR             0
        GBP_MYR             0
        EUR_MYR             0
        JPY100_MYR          0
        CHF_MYR             0
        AUD_MYR             0
        CAD_MYR             0
        SGD_MYR             0
        HKD100_MYR          0
        dtype: int64
```

```
In [9]: 1 df.dtypes
```

```
Out[9]: Date                object
        USD_MYR             object
        GBP_MYR             object
        EUR_MYR             object
        JPY100_MYR          object
        CHF_MYR             object
        AUD_MYR             object
        CAD_MYR             object
        SGD_MYR             object
        HKD100_MYR          object
        dtype: object
```

```
In [10]: 1 #B. Data Visualisation
```

```
In [11]: 1 for col in df.columns[1:]:
        2     df[col] = pd.to_numeric(df[col], errors='coerce')
```

```
In [12]: 1 df['Date'] = pd.to_datetime(df['Date'])
        2 df['month'] = df['Date'].dt.month
        3 df['year'] = df['Date'].dt.year
        4 df['month_year'] = df['Date'].dt.to_period('M')
```


```
In [13]: 1 df_groupby_myr = df.groupby('month_year').USD_MYR.mean().reset_index()  
2 df_groupby_myr
```

Out[13]:

	month_year	USD_MYR
0	2000-01	4.659383
1	2000-02	4.623653
2	2000-03	4.608377
3	2000-04	4.603860
4	2000-05	4.521210
...
271	2022-08	5.650715
272	2022-09	5.609749
273	2022-10	5.696011
274	2022-11	5.770798
275	2022-12	5.706441

276 rows × 2 columns

```
In [14]: 1 from matplotlib import pyplot as plt
2
3 x = df_groupby_myr['month_year'].astype(str)
4 y = df_groupby_myr['USD_MYR']
5
6 plt.figure(figsize=(8,4))
7 plt.plot(x, y)
8 plt.title("Exchange Rate: USD/MYR")
9 plt.xlabel("Month")
10 plt.ylabel("Exchange Rate")
11 plt.show()
```



```
In [15]: 1 df.columns
```

```
Out[15]: Index(['Date', 'USD_MYR', 'GBP_MYR', 'EUR_MYR', 'JPY100_MYR', 'CHF_MYR',
               'AUD_MYR', 'CAD_MYR', 'SGD_MYR', 'HKD100_MYR', 'month', 'year',
               'month_year'],
              dtype='object')
```

```

In [16]: 1 groupby_all = df.groupby('month_year').apply(lambda x: pd.Series({
2                                                    'USD_mean': x['USD_MYR']
3                                                    'GBP_mean': x['GBP_MYR']
4                                                    'EUR_mean': x['EUR_MYR']
5                                                    'JPY_mean': x['JPY100_MY
6                                                    'CHF_mean': x['CHF_MYR']
7                                                    'AUD_mean': x['AUD_MYR']
8                                                    'CAD_mean': x['CAD_MYR']
9                                                    'SGD_mean': x['SGD_MYR']
10                                                   'HKD_mean': x['HKD100_MY
11 })).reset_index()
12
13 groupby_all = pd.DataFrame(groupby_all)
14
15 groupby_all['MYR_mean'] = [1 for i in range(276)]
16 groupby_all.head()

```


Out[16]:

	month_year	USD_mean	GBP_mean	EUR_mean	JPY_mean	CHF_mean	AUD_mean	CAD_mean
0	2000-01	4.659383	5.209283	5.394987	1.318802	0.815958	34.607715	2.611106
1	2000-02	4.623653	5.150284	5.372135	1.268332	0.793151	34.569567	2.576541
2	2000-03	4.608377	5.096289	5.349577	1.300361	0.779052	34.547047	2.568351
3	2000-04	4.603860	5.079511	5.354026	1.314547	0.778458	34.528688	2.558911
4	2000-05	4.521210	4.938006	5.269411	1.284524	0.753359	34.507037	2.507741

```

In [17]: 1 y_all = groupby_all[['USD_mean', 'GBP_mean', 'EUR_mean', 'JPY_mean', 'CHF_me
2 labels = ["USD_mean", "GBP_mean", "EUR_mean", "JPY_mean", "CHF_mean", "AUD_m
3 x_ticks = list(range(1, 280, 13))
4 x_ticklabels = [x for x in range(2000, 2022)]
5
6 plt.figure(figsize=(20,15))
7 ax = plt.subplot()
8
9 plt.plot(x, y_all)
10 ax.set_xticks(x_ticks)
11 ax.set_xticklabels(x_ticklabels)
12 plt.legend(labels)
13 plt.title("Exchange Rate: Top Countries/MYR")
14 plt.xlabel("Year")
15 plt.ylabel("Exchange Rate")
16
17 plt.show()


```



In [18]:

```
1 from bokeh.models import ColumnDataSource, HoverTool
2 from bokeh.models.annotations import Title
3 from bokeh.plotting import figure, show
4 from bokeh.io import output_notebook
5 from bokeh.palettes import Spectral10
6
7 from datetime import datetime
8
9 groupby_all['month_year'] = groupby_all['month_year'].astype(str)
10 groupby_all['month_year'] = pd.to_datetime(groupby_all['month_year'])
11
12 source = ColumnDataSource(groupby_all)
13
14 output_notebook()
15
16 p = figure(plot_height = 500, \
17           plot_width = 900, \
18           x_axis_type = 'datetime')
19
20 p.line(x='month_year', y='USD_mean', legend_label = 'USD_mean', source=source)
21 p.line(x='month_year', y='GBP_mean', legend_label = 'GBP_mean', source=source)
22 p.line(x='month_year', y='EUR_mean', legend_label = 'EUR_mean', source=source)
23 p.line(x='month_year', y='JPY_mean', legend_label = 'JPY_mean', source=source)
24 p.line(x='month_year', y='CHF_mean', legend_label = 'CHF_mean', source=source)
25 p.line(x='month_year', y='AUD_mean', legend_label = 'AUD_mean', source=source)
26 p.line(x='month_year', y='CAD_mean', legend_label = 'CAD_mean', source=source)
27 p.line(x='month_year', y='SGD_mean', legend_label = 'SGD_mean', source=source)
28 p.line(x='month_year', y='HKD_mean', legend_label = 'HKD_mean', source=source)
29 p.line(x='month_year', y='MYR_mean', legend_label = 'MYR_mean', source=source)
30
31
32 p.xaxis[0].ticker.desired_num_ticks = 10
33
34 p.xaxis.axis_label = 'Year'
35 p.yaxis.axis_label = 'Exchange Rate'
36
37 t = Title()
38 t.text = 'Exchange Rate: Top Countries/MYR'
39 p.title = t
40
41 hover = HoverTool(tooltips = [
42     ('USD_mean', '@USD_mean'),\
43     ('GBP_mean', '@GBP_mean'),\
44     ('EUR_mean', '@EUR_mean'),\
45     ('JPY_mean', '@JPY_mean'),\
46     ('CHF_mean', '@CHF_mean'),\
47     ('AUD_mean', '@AUD_mean'),\
48     ('CAD_mean', '@CAD_mean'),\
49     ('SGD_mean', '@SGD_mean'),\
50     ('HKD_mean', '@HKD_mean'),\
51     ('MYR_mean', '@MYR_mean')])
52 p.add_tools(hover)
53 show(p)
```

(<http://bokeh.pydata.org/en/latest/docs/0.12.0/quickstart.html>) BokehJS 2.4.3 successfully loaded.



In [19]: 1 *#C. Hypothesis Testing*

In [20]: 1 `usd_aud = df.loc[:, ['USD_MYR', 'AUD_MYR']]`
 2 `print(usd_aud)`

	USD_MYR	AUD_MYR
0	3.8000	2.5002
1	3.8000	2.5021
2	3.8000	2.4915
3	3.8000	2.5002
4	3.8000	2.4854
...
17008	0.0187	120.5690
17009	0.0188	120.3763
17010	0.0188	120.6643
17011	0.0187	120.3169
17012	0.0187	120.1601

[17013 rows x 2 columns]

```
In [21]: 1  usd = np.array(usd_aud.iloc[:,0])
2  aud = np.array(usd_aud.iloc[:,1])
3
4  import statsmodels.stats.power as sms
5  from statsmodels import stats
6
7  #find sample size through power analysis
8  n = sms.TTestPower().solve_power(0.3, power=0.9, alpha=0.05)
9  print(n)
```

118.68650942951783

```
In [22]: 1  #generate sample sets from population
2  aud_sample = np.random.choice(aud, size=118)
3  usd_sample = np.random.choice(usd, size=118)
4
5  print(np.array(list(zip(aud_sample, usd_sample))))
```

```
[9.039160e+01 3.140000e+00]
[9.084940e+01 1.011290e+01]
[8.684750e+01 4.475000e+00]
[1.196355e+02 3.800000e+00]
[0.000000e+00 1.037790e+01]
[9.371540e+01 2.160000e-02]
[0.000000e+00 3.800000e+00]
[2.157900e+00 3.779200e+00]
[2.651100e+00 3.800000e+00]
[8.166010e+01 4.090500e+00]
[1.069396e+02 3.028100e+00]
[8.607240e+01 1.034320e+01]
[3.004800e+00 2.140000e-02]
[9.359580e+01 2.110000e-02]
[2.928200e+00 1.810000e-02]
[8.249720e+01 0.000000e+00]
[2.819800e+00 9.991400e+00]
[2.961600e+00 1.810000e-02]
[2.888800e+00 1.009160e+01]
[0.102440e+01 0.000000e+00]
```

```
In [23]: 1  from scipy.stats import shapiro, pearsonr
2
3  #test normality of each sample set (H0: normally distributed)
4  stat, p1 = shapiro(aud_sample)
5  stat, p2 = shapiro(usd_sample)
6  print(p1)
7  print(p2)
8
9  #test whether the sample sets are correlated or not (H0: not correlated)
10 stat, p3 = pearsonr(aud_sample, usd_sample)
11 print(p3)
```

```
7.05752622141842e-13
2.89434587408266e-09
0.08451891450627709
```

```
In [24]: 1 #scipy
2 from scipy.stats import ttest_ind
3 tstat, pval = ttest_ind(aud_sample, usd_sample)
4 print('%.10f' %pval)
5
6 #statsmodels
7 import statsmodels.stats.api as sm
8 tstat, pval2, df = sm.ttest_ind(aud_sample, usd_sample)
9 print('%.30f' %pval2)
```

```
0.0000000000
0.000000000000000000091361534238
```

```
In [25]: 1 #D. Machine Learning Models
```

```
In [26]: 1 ir_df = pd.read_csv("interbank-intraday-rates-2018-2022.csv")
2 ir_df
```

Out[26]:

	Trading_Date	Highest_Rate	Lowest_Rate
0	2-Jan-18	4.045	4.0175
1	3-Jan-18	4.028	4.0140
2	4-Jan-18	4.028	4.0060
3	5-Jan-18	4.005	3.9900
4	8-Jan-18	3.999	3.9860
...
1189	21-Dec-22	4.441	4.4350
1190	22-Dec-22	4.435	4.4220
1191	23-Dec-22	4.435	4.4200
1192	28-Dec-22	4.440	4.4180
1193	29-Dec-22	4.429	4.4200

1194 rows × 3 columns

```
In [27]: 1 ir_df.dtypes
```

```
Out[27]: Trading_Date      object
Highest_Rate      float64
Lowest_Rate      float64
dtype: object
```

```
In [28]: 1 for col in ir_df.columns[1:]:
2     ir_df[col] = pd.to_numeric(ir_df[col], errors='coerce')
```

```
In [29]: 1 ir_df['Trading_Date'] = pd.to_datetime(ir_df['Trading_Date'])
2 ir_df['month'] = ir_df['Trading_Date'].dt.month
3 ir_df['year'] = ir_df['Trading_Date'].dt.year
4 ir_df['month_year'] = ir_df['Trading_Date'].dt.to_period('M')
```

In [30]: 1 ir_df

Out[30]:

	Trading_Date	Highest_Rate	Lowest_Rate	month	year	month_year
0	2018-01-02	4.045	4.0175	1	2018	2018-01
1	2018-01-03	4.028	4.0140	1	2018	2018-01
2	2018-01-04	4.028	4.0060	1	2018	2018-01
3	2018-01-05	4.005	3.9900	1	2018	2018-01
4	2018-01-08	3.999	3.9860	1	2018	2018-01
...
1189	2022-12-21	4.441	4.4350	12	2022	2022-12
1190	2022-12-22	4.435	4.4220	12	2022	2022-12
1191	2022-12-23	4.435	4.4200	12	2022	2022-12
1192	2022-12-28	4.440	4.4180	12	2022	2022-12
1193	2022-12-29	4.429	4.4200	12	2022	2022-12

In [31]: 1 df_groupby_ir = ir_df.groupby('month_year').Highest_Rate.mean().reset_index()
2 df_groupby_ir.head()

Out[31]:

	month_year	Highest_Rate
0	2018-01	3.966976
1	2018-02	3.918233
2	2018-03	3.908182
3	2018-04	3.889775
4	2018-05	3.969250

In [32]: 1 x_ir = df_groupby_ir['Highest_Rate']
2 x_ir = np.array(x_ir).reshape(-1,1)

In [33]: 1 usd_myr_fx = df_groupby_myr[(df_groupby_myr['month_year'] >= '2018-01') & (d
2 usd_myr = usd_myr_fx['USD_MYR']
3
4 y_fx = usd_myr

In [34]: 1 from sklearn.model_selection import train_test_split
2 from sklearn.linear_model import LinearRegression

In [35]: 1 x_train, x_test, y_train, y_test = train_test_split(x_ir, y_fx, train_size=0

Model to be used in this dataet

1. Linear Regression
2. Polynomial Regression

3. Ridge Regression & Lasso Regression
4. K-Nearest Neighbour Regression
5. Support Vector Regression

```
In [36]: 1 R_squared=[]  
2 MAE=[]  
3 MSE=[]  
4 RMSE=[]  
5 Model = []
```


Linear Regression

```
In [37]: 1 model = LinearRegression()  
2 model.fit(x_train, y_train)  
3 y_predict = model.predict(x_test)  
4 print(y_predict)  
  
[5.70250591 5.705314    5.64749862 5.7190307  5.722121    5.71360114  
 5.70990829 5.68609478 5.70389803 5.76175054 5.69188168 5.77177564]
```

```
In [38]: 1 print("R-squared of training data is: " + str(model.score(x_train, y_train)))  
2 print("R-squared of testing data is: " + str(model.score(x_test, y_test)))  
  
R-squared of training data is: 0.06094924772473809  
R-squared of testing data is: 0.06983088859191167
```

```
In [39]: 1 model = LinearRegression()  
2 model.fit(x_ir, y_fx)  
3 y_fx_predict = model.predict(x_ir)  
4  
5 print("R-squared of the entire dataset is: " + str(model.score(x_ir, y_fx)))  
  
R-squared of the entire dataset is: 0.06231308561643811
```

```
In [40]: 1 month_year = usd_myr_fx['month_year']
2 month_year = month_year.astype(str)
3
4 plt.figure(figsize=(12,6))
5 plt.scatter(month_year, y_fx, alpha=0.4)
6 plt.plot(month_year, y_fx_predict)
7 plt.title("Linear Regression: AUD/USD Exchange Rate (1 var: Interest Rate)")
8 plt.xlabel("Month-Year")
9 plt.ylabel("Exchange Rate")
10 plt.xticks(fontsize=4)
11 plt.show()
```



```
In [41]: 1 import sklearn.metrics as metrics
2
3 mae_linear = metrics.mean_absolute_error(y_fx, y_fx_predict)
4 mse_linear = metrics.mean_squared_error(y_fx, y_fx_predict)
5 rmse_linear = np.sqrt(mse_linear) # or mse**(0.5)
6 r2_linear = metrics.r2_score(y_fx, y_fx_predict)
7
8 print("Results of sklearn.metrics: Linear")
9 print("MAE:", mae_linear)
10 print("MSE:", mse_linear)
11 print("RMSE:", rmse_linear)
12 print("R-Squared:", r2_linear)
```

```
Results of sklearn.metrics: Linear
MAE: 0.13781033888268673
MSE: 0.024034915726698233
RMSE: 0.15503198291545597
R-Squared: 0.06231308561643811
```


```
In [42]: 1 MAE.append(mae_linear)
2 MSE.append(mse_linear)
3 RMSE.append(rmse_linear)
4 R_squared.append(r2_linear)
5 Model.append('Linear')
```

Polynomial Regression

```
In [43]: 1 from sklearn.preprocessing import PolynomialFeatures
2 poly = PolynomialFeatures(degree=4)
3 x_poly = poly.fit_transform(x_ir)
4 model_poly = LinearRegression()
5 model_poly.fit(x_poly, y_fx)
6 y_pred = model_poly.predict(x_poly)
7 print(model_poly.score(x_poly, y_fx))
```

0.30482494556026196

```
In [44]: 1 month_year = usd_myr_fx['month_year']
2 month_year = month_year.astype(str)
3
4 plt.figure(figsize=(12,6))
5 plt.scatter(month_year, y_fx, alpha=0.4)
6 plt.plot(month_year, y_pred)
7 plt.title("Polynomial Regression: AUD/USD Exchange Rate (1 var: Interest Rat
8 plt.xlabel("Month-Year")
9 plt.ylabel("Exchange Rate")
10 plt.xticks(fontsize=4)
11 plt.show()
```



```
In [45]: 1 import sklearn.metrics as metrics
2
3 mae_poly = metrics.mean_absolute_error(y_fx, y_pred)
4 mse_poly = metrics.mean_squared_error(y_fx, y_pred)
5 rmse_poly = np.sqrt(mse_poly) # or mse**(0.5)
6 r2_poly = metrics.r2_score(y_fx, y_pred)
7
8 print("Results of sklearn.metrics: Polynomial")
9 print("MAE:", mae_poly)
10 print("MSE:", mse_poly)
11 print("RMSE:", rmse_poly)
12 print("R-Squared:", r2_poly)
```

```
Results of sklearn.metrics: Polynomial
MAE: 0.11444499943671242
MSE: 0.0178188194721115
RMSE: 0.13348715096259828
R-Squared: 0.30482494556026196
```

```
In [46]: 1 MAE.append(mae_poly)
2 MSE.append(mse_poly)
3 RMSE.append(rmse_poly)
4 R_squared.append(r2_poly)
5 Model.append('Polynomial')
```

Ridge Regression & Lasso Regression



```

In [47]: 1 from sklearn.linear_model import Ridge, Lasso
2
3 ridge = Ridge(alpha=0.001)
4 ridge.fit(x_ir, y_fx)
5
6 y_fx_ridge = ridge.predict(x_ir)
7 print(ridge.score(x_ir, y_fx))
8
9 lasso = Lasso(alpha=0.001)
10 lasso.fit(x_ir, y_fx)
11
12 y_fx_lasso = lasso.predict(x_ir)
13 print(lasso.score(x_ir, y_fx))
14
15 labels = ["y_entire", "Ridge", "Lasso"]
16
17 plt.figure(figsize=(14,5))
18 plt.scatter(month_year, y_fx, alpha=0.4)
19 plt.plot(month_year, y_fx_ridge, color='r')
20 plt.plot(month_year, y_fx_lasso, color='b')
21 plt.legend(labels)
22 plt.title("Ridge & Lasso Regression: USD/MYR Exchange Rate ")
23 plt.xlabel("Month-Year")
24 plt.ylabel("Exchange Rate")
25 plt.xticks(fontsize=5)
26 plt.show()

```

0.062313058289344525

0.060761908867496994



```
In [48]: 1 import sklearn.metrics as metrics
2
3 mae_ridge = metrics.mean_absolute_error(y_fx, y_fx_ridge)
4 mse_ridge = metrics.mean_squared_error(y_fx, y_fx_ridge)
5 rmse_ridge = np.sqrt(mse_ridge) # or mse**(0.5)
6 r2_ridge = metrics.r2_score(y_fx, y_fx_ridge)
7
8 print("Results of sklearn.metrics: Ridge")
9 print("MAE:", mae_ridge)
10 print("MSE:", mse_ridge)
11 print("RMSE:", rmse_ridge)
12 print("R-Squared:", r2_ridge)
```

Results of sklearn.metrics: Ridge
MAE: 0.13780718221029928
MSE: 0.024034916427149933
RMSE: 0.15503198517451144
R-Squared: 0.062313058289344525

```
In [49]: 1 MAE.append(mae_ridge)
2 MSE.append(mse_ridge)
3 RMSE.append(rmse_ridge)
4 R_squared.append(r2_ridge)
5 Model.append('Ridge')
```

```
In [50]: 1 import sklearn.metrics as metrics
2
3 mae_lasso = metrics.mean_absolute_error(y_fx, y_fx_lasso)
4 mse_lasso = metrics.mean_squared_error(y_fx, y_fx_lasso)
5 rmse_lasso = np.sqrt(mse_lasso) # or mse**(0.5)
6 r2_lasso = metrics.r2_score(y_fx, y_fx_lasso)
7
8 print("Results of sklearn.metrics: Lasso")
9 print("MAE:", mae_lasso)
10 print("MSE:", mse_lasso)
11 print("RMSE:", rmse_lasso)
12 print("R-Squared:", r2_lasso)
```

Results of sklearn.metrics: Lasso
MAE: 0.13705825992359108
MSE: 0.024074675695474723
RMSE: 0.15516016143158243
R-Squared: 0.060761908867496994

```
In [51]: 1 MAE.append(mae_lasso)
2 MSE.append(mse_lasso)
3 RMSE.append(rmse_lasso)
4 R_squared.append(r2_lasso)
5 Model.append('Lasso')
```

K-Nearest Neighbour Regression

```
In [52]: 1 from sklearn.neighbors import KNeighborsRegressor
2 knn_model = KNeighborsRegressor(n_neighbors=5, metric='euclidean')
3 knn_model.fit(x_train, y_train)
4 knn_pred = knn_model.predict(x_test)
```

```
In [53]: 1 print("R-squared of training data is: " + str(knn_model.score(x_train, y_train)))
2 print("R-squared of testing data is: " + str(knn_model.score(x_test, y_test)))
```


R-squared of training data is: 0.426644008702908
R-squared of testing data is: 0.10020811407667907

```
In [54]: 1 model.fit(x_ir, y_fx)
2 y_knn = knn_model.predict(x_ir)
```

```
In [55]: 1 print("R-squared of the entire dataset is: " + str(knn_model.score(x_ir, y_fx)))
```

R-squared of the entire dataset is: 0.38673773790801846

```
In [56]: 1 month_year = usd_myr_fx['month_year']
2 month_year = month_year.astype(str)
3
4 plt.figure(figsize=(12,6))
5 plt.scatter(month_year, y_fx, alpha=0.4)
6 plt.plot(month_year, y_knn)
7 plt.title("kNN Regression: AUD/USD Exchange Rate (1 var: Interest Rate)")
8 plt.xlabel("Month-Year")
9 plt.ylabel("Exchange Rate")
10 plt.xticks(fontsize=4)
11 plt.show()
```



```
In [57]: 1 import sklearn.metrics as metrics
2
3 mae_knn = metrics.mean_absolute_error(y_fx, y_knn)
4 mse_knn = metrics.mean_squared_error(y_fx, y_knn)
5 rmse_knn = np.sqrt(mse_knn) # or mse**(0.5)
6 r2_knn = metrics.r2_score(y_fx, y_knn)
7
8 print("Results of sklearn.metrics: kNN")
9 print("MAE:", mae_knn)
10 print("MSE:", mse_knn)
11 print("RMSE:", rmse_knn)
12 print("R-Squared:", r2_knn)
```

Results of sklearn.metrics: kNN
MAE: 0.10378981723507695
MSE: 0.015719219882080818
RMSE: 0.12537631308218
R-Squared: 0.38673773790801846

```
In [58]: 1 MAE.append(mae_knn)
2 MSE.append(mse_knn)
3 RMSE.append(rmse_knn)
4 R_squared.append(r2_knn)
5 Model.append('KNN')
```

Support Vector Regression

```
In [59]: 1 from sklearn.svm import SVR
2 svr_model = SVR(kernel = 'rbf')
3 svr_model.fit(x_train, y_train)
4 svr_pred = svr_model.predict(x_test)
5 print(svr_pred)
```

[5.73497741 5.73975219 5.44023252 5.76442358 5.77444343 5.75203555
5.7462519 5.6650441 5.73748163 5.84743363 5.69983773 5.75582786]


```
In [60]: 1 print("R-squared of training data is: " + str(svr_model.score(x_train, y_train)))
2 print("R-squared of testing data is: " + str(svr_model.score(x_test, y_test)))
```

R-squared of training data is: 0.39299413096649005
R-squared of testing data is: 0.2076951343642901

```
In [61]: 1 svr_model = SVR(kernel = 'rbf')
2 svr_model.fit(x_ir, y_fx)
3 y_svr = svr_model.predict(x_ir)
4
5 print("R-squared of the entire dataset is: " + str(svr_model.score(x_ir, y_fx)))
```

R-squared of the entire dataset is: 0.355006238035672

```
In [62]: 1 month_year = usd_myr_fx['month_year']
2 month_year = month_year.astype(str)
3
4 plt.figure(figsize=(12,6))
5 plt.scatter(month_year, y_fx, alpha=0.4)
6 plt.plot(month_year, y_svr)
7 plt.title("Linear Regression: AUD/USD Exchange Rate (1 var: Interest Rate)")
8 plt.xlabel("Month-Year")
9 plt.ylabel("Exchange Rate")
10 plt.xticks(fontsize=4)
11 plt.show()
```




```
In [63]: 1 import sklearn.metrics as metrics
2
3 mae_svr = metrics.mean_absolute_error(y_fx, y_svr)
4 mse_svr = metrics.mean_squared_error(y_fx, y_svr)
5 rmse_svr = np.sqrt(mse_svr) # or mse**(0.5)
6 r2_svr = metrics.r2_score(y_fx, y_svr)
7
8 print("Results of sklearn.metrics: SVR")
9 print("MAE:", mae_svr)
10 print("MSE:", mse_svr)
11 print("RMSE:", rmse_svr)
12 print("R-Squared:", r2_svr)
```

Results of sklearn.metrics: SVR
MAE: 0.10831126565444975
MSE: 0.016532565907939525
RMSE: 0.1285790259254577
R-Squared: 0.355006238035672

```
In [64]: 1 MAE.append(mae_svr)
2 MSE.append(mse_svr)
3 RMSE.append(rmse_svr)
4 R_squared.append(r2_svr)
5 Model.append('SVR')
```

Result

```
In [65]: 1 labels = ["y_entire", "Polynomial", "Ridge", "Lasso", "KNN", "SVR"]
2
3 plt.figure(figsize=(14,5))
4 plt.scatter(month_year, y_fx, alpha=0.4)
5 plt.plot(month_year, y_pred, color='black')
6 plt.plot(month_year, y_fx_ridge, color='r')
7 plt.plot(month_year, y_fx_lasso, color='b')
8 plt.plot(month_year, y_knn, color='y')
9 plt.plot(month_year, y_svr, color='g')
10 plt.legend(labels)
11 plt.title("Result: USD/MYR Exchange Rate ")
12 plt.xlabel("Month-Year")
13 plt.ylabel("Exchange Rate")
14 plt.xticks(fontsize=5)
15 plt.show()
```



Evaluation

```
In [66]: 1 print("Model Used:-")
2 Model
```

Model Used:-

```
Out[66]: ['Linear', 'Polynomial', 'Ridge', 'Lasso', 'KNN', 'SVR']
```

```
In [67]: 1 output=pd.DataFrame({'Model':Model,
2                        'MAE':[mae_linear, mae_poly,mae_ridge, mae_lasso, mae_kn
3                        'MSE':[mse_linear, mse_poly ,mse_ridge, mse_lasso, mse_k
4                        'RMSE': [rmse_linear, rmse_poly ,rmse_ridge, rmse_lasso
5                        'R_squared':[r2_linear, r2_poly , r2_ridge, r2_lasso, r
6                        ]})
7 output=output.set_index('Model')
8 output
```

Out[67]:

	MAE	MSE	RMSE	R_squared
Model				
Linear	0.137810	0.024035	0.155032	0.062313
Polynomial	0.114445	0.017819	0.133487	0.304825
Ridge	0.137807	0.024035	0.155032	0.062313
Lasso	0.137058	0.024075	0.155160	0.060762
KNN	0.103790	0.015719	0.125376	0.386738
SVR	0.108311	0.016533	0.128579	0.355006

Based on the result that we have run, we conclude that KNN is the best model among other model. This is because MAE, MSE, RMSE of KNN is the lowest, which is 0.103790, 0.015719, 0.125376, 0.386738 while the R-Squared is the highest which is 0.386738. In addition, the best evaluation is RMSE due to the time series forecasting and measure the average difference between the predicted and actual values.

Furthermore, MAE also another option to evaluate the exchange rate. Both evaluations give a good indication of how accurate the model in predicting exchange rates.(in applied predictive modelling,Max Kuhn, J.Kjell)